

Case Study Brief

Hello! Thank you for applying with us as a backend developer. This mini project **should be completed within 5 days after you have received this document**. Please spare your time to complete this project with the best results. We are really pleased to answer your questions if there are unclear things.

Objective

Your mission is to build a backend service that automates the initial screening of a job application. The service will receive a candidate's CV and a project report, evaluate them against a specific job description and a case study brief, and produce a structured, AI-generated evaluation report.

Core Logic & Data Flow

The system operates with a clear separation of inputs and reference documents:

Candidate-Provided Inputs (The Data to be Evaluated):

- Candidate CV:** The candidate's resume (PDF).
- Project Report:** The candidate's project report to our take-home case study (PDF)

System-Internal Documents (The "Ground Truth" for Comparison):

- Job Description:** A document detailing the requirements and responsibilities for the role — You can use the job description you're currently applying. This document will be used as ground truth for **Candidate CV**.
 - To make sure the vector retrieval is accurate enough, you might need to ingest a few job description documents as well.
- Case Study Brief:** This document. Used as ground truth for **Project Report**. (PDF)
- Scoring Rubric:** A predefined set of parameters for evaluating CV and Report, each has it's own documents. (PDF)

We want to see your ability to combine **backend engineering** with **AI workflows** (prompt design, LLM chaining, retrieval, resilience).

Deliverables

1. Backend Service (API endpoints)

Implement a backend service with at least the following RESTful API endpoints:

- `POST /upload`
 - Accepts multipart/form-data containing the Candidate CV and Project Report (PDF).
 - Stores these files, return each with it's own ID for later processing.
- `POST /evaluate`
 - Triggers the asynchronous AI evaluation pipeline. Receives input job title (string), and **both** document ID.
 - Immediately returns a job ID to track the evaluation process.

```
{
  "id": "456",
  "status": "queued"
}
```

- `GET /result/{id}`
 - Retrieves the status and result of an evaluation job. This endpoint should reflect the asynchronous, **multi-stage** nature of the process.
 - Possible responses:
 - While queued or processing

```
{
  "id": "456",
  "status": "queued" | "processing"
}
```

- Once completed

```
{
  "id": "456",
  "status": "completed",
  "result": {
    "cv_match_rate": 0.82,
    "cv_feedback": "Strong in backend and cloud, limited AI integration experience...",
    "project_score": 4.5,
    "project_feedback": "Meets prompt chaining requirements, lacks error handling robustness...",
    "overall_summary": "Good candidate fit, would benefit from deeper RAG knowledge..."
  }
}
```

2. Evaluation Pipeline

Design and implement an AI-driven pipeline which will be triggered by **[POST] /evaluate** endpoint. Should consist these key Components:

- **RAG (Context Retrieval)**
 - Ingest all **System-Internal Documents** (Job Description, Case Study Brief, Both Scoring Rubrics) into a vector database.
 - Retrieve relevant sections and inject into prompts (e.g., “for CV scoring” vs “for project scoring”).

- **Prompt Design & LLM Chaining**

The pipeline should consists of

- CV Evaluation
 - Parse the candidate's CV into structured data.
 - Retrieve relevant information from both Job Description and CV Scoring Rubrics.
 - Use an LLM to get these result: `cv_match_rate` & `cv_feedback`
- Project Report Evaluation
 - Parse the candidate's Project Report into structured data.
 - Retrieve relevant information from both Case Study Brief and CV Scoring Rubrics.
 - Use an LLM to get these result: `project_score` & `project_feedback`
- Final Analysis
 - Use a final LLM call to synthesize the outputs from previous steps into a concise `overall_summary`.
- **Long-Running Process Handling**
 - `POST /evaluate` should **not block** until LLM Chaining finishes.
 - Store task, return job ID, allow `GET /result/{id}` to check later periodically.
- **Error Handling & Randomness Control**
 - Simulate any edge cases you can think of and how well your service can handle them.
 - Simulate failures from LLM API (timeouts, rate limit).
 - Implement retries/back-off.
 - Control LLM temperature or add validation layer to keep responses stable.

3. Standardized Evaluation Parameters

Define at least these scoring parameters:

CV Evaluation (Match Rate)

- **Technical Skills Match** (backend, databases, APIs, cloud, AI/LLM exposure).
- **Experience Level** (years, project complexity).
- **Relevant Achievements** (impact, scale).
- **Cultural Fit** (communication, learning attitude).

Project Deliverable Evaluation

- **Correctness** (meets requirements: prompt design, chaining, RAG, handling errors).
- **Code Quality** (clean, modular, testable).
- **Resilience** (handles failures, retries).

- **Documentation** (clear README, explanation of trade-offs).
- **Creativity / Bonus** (optional improvements like authentication, deployment, dashboards).

Each parameter can be scored **1–5**, then aggregated to final score.

Requirements

- Use any backend framework (Rails, Django, Node.js, etc.).
- Use a proper LLM service (e.g., OpenAI, Gemini, or OpenRouter). There are several free LLM API providers available.
- Use a simple **vector DB** (e.g. ChromaDB, Qdrant, etc) or **RAG-as-a-service** (e.g. Ragie, S3 Vector, etc), any of your own choice.
- Provide README with run instructions + explanation of design choices.
- Provide the documents together with their ingestion scripts in the repository for reproducability purposes.

Scoring Rubric for Case Study Evaluation

CV Match Evaluation (1–5 scale per parameter)

Parameter	Description	Scoring Guide
Technical Skills Match (Weight: 40%)	Alignment with job requirements (backend, databases, APIs, cloud, AI/LLM).	1 = Irrelevant skills, 2 = Few overlaps, 3 = Partial match, 4 = Strong match, 5 = Excellent match + AI/LLM exposure
Experience Level (Weight: 25%)	Years of experience and project complexity.	1 = <1 yr / trivial projects, 2 = 1–2 yrs, 3 = 2–3 yrs with mid-scale projects, 4 = 3–4 yrs solid track record, 5 = 5+ yrs / high-impact projects
Relevant Achievements (Weight: 20%)	Impact of past work (scaling, performance, adoption).	1 = No clear achievements, 2 = Minimal improvements, 3 = Some measurable outcomes, 4 = Significant contributions, 5 = Major measurable impact
Cultural / Collaboration Fit (Weight: 15%)	Communication, learning mindset, teamwork/leadership.	1 = Not demonstrated, 2 = Minimal, 3 = Average, 4 = Good, 5 = Excellent and well-demonstrated

Project Deliverable Evaluation (1–5 scale per parameter)

Parameter	Description	Scoring Guide
Correctness (Prompt & Chaining) (Weight: 30%)	Implements prompt design, LLM chaining, RAG context injection.	1 = Not implemented, 2 = Minimal attempt, 3 = Works partially, 4 = Works correctly, 5 = Fully correct + thoughtful
Code Quality & Structure (Weight: 25%)	Clean, modular, reusable, tested.	1 = Poor, 2 = Some structure, 3 = Decent modularity, 4 = Good structure + some tests, 5 = Excellent quality + strong tests
Resilience & Error Handling (Weight: 20%)	Handles long jobs, retries, randomness, API failures.	1 = Missing, 2 = Minimal, 3 = Partial handling, 4 = Solid handling, 5 = Robust, production-ready
Documentation & Explanation (Weight: 15%)	README clarity, setup instructions, trade-off explanations.	1 = Missing, 2 = Minimal, 3 = Adequate, 4 = Clear, 5 = Excellent + insightful
Creativity / Bonus (Weight: 10%)	Extra features beyond requirements.	1 = None, 2 = Very basic, 3 = Useful extras, 4 = Strong enhancements, 5 = Outstanding creativity

3. Overall Candidate Evaluation

- CV Match Rate:** Weighted Average (1–5) → Convert to 0-1 decimal (×0.2).
- Project Score:** Weighted Average (1–5)
- Overall Summary:** Service should return 3–5 sentences (strengths, gaps, recommendations).

Study Case Submission Template

Please use this template to document your solution. Submit it as a **PDF file** along with your project repository.

1. Title

2. Candidate Information

- **Full Name:**
- **Email Address:**


3. Repository Link

- Provide a link to your GitHub repository.
- **▲ Important:** Do **not** use the word *Rakamin* anywhere in your repository name, commits, or documentation. This is to reduce plagiarism risk.
- Example: *github.com/username/ai-cv-evaluator*

4. Approach & Design (Main Section)

Tell the story of how you approached this challenge. We want to understand your thinking process, not just the code. Please include:

- **Initial Plan**
 - How you broke down the requirements.
 - Key assumptions or scope boundaries.
- **System & Database Design**
 - API endpoints design.
 - Database schema (diagram or explanation).
 - Job queue / long-running task handling.
- **LLM Integration**
 - Why you chose a specific LLM or provider.
 - Prompt design decisions.
 - Chaining logic (if any).
 - RAG (retrieval, embeddings, vector DB) strategy.
- **Prompting Strategy** (examples of your actual prompts)
- **Resilience & Error Handling**
 - How you handled API failures, timeouts, or randomness.
 - Any retry, backoff, or fallback logic.
- **Edge Cases Considered**
 - What unusual inputs or scenarios you thought about.
 - How you tested them.

 This is your chance to be a storyteller. Imagine you're presenting to a CTO, clarity and reasoning matter more than buzzwords.

5. Results & Reflection

- **Outcome**
 - What worked well in your implementation?
 - What didn't work as expected?
- **Evaluation of Results**
 - If the evaluation scores/outputs were bad or inconsistent, explain why.
 - If they were good, explain what made them stable.
- **Future Improvements**
 - What would you do differently with more time?
 - What constraints (time, tools, API limits) affected your solution?

6. Screenshots of Real Responses

- Show **real JSON response** from your API using your own **CV + Project Report**.
- Minimum:
 - `/evaluate` → returns job_id + status
 - `/result/:id` → returns final evaluation (scores + feedback)
- Paste screenshots or Postman/terminal logs.

7. (Optional) Bonus Work

If you added extra features, describe them here.

Job Description - Product Engineer (Backend) 2025

Rakamin is hiring a Product Engineer (Backend) to work on Rakamin. We're looking for dedicated engineers who write code they're proud of and who are eager to keep scaling and improving complex systems, including those powered by AI.

About the Job

You'll be building new product features alongside a frontend engineer and product manager using our Agile methodology, as well as addressing issues to ensure our apps are robust and our codebase is clean. As a Product Engineer, you'll write clean, efficient code to enhance our product's codebase in meaningful ways.

In addition to classic backend work, this role also touches on building AI-powered systems, where you'll design and orchestrate how large language models (LLMs) integrate into Rakamin's product ecosystem.

Here are some real examples of the work in our team:

- Collaborating with frontend engineers and 3rd parties to build robust backend solutions that support highly configurable platforms and cross-platform integration.
- Developing and maintaining server-side logic for central database, ensuring high performance throughput and response time.
- Designing and fine-tuning AI prompts that align with product requirements and user contexts.
- Building LLM chaining flows, where the output from one model is reliably passed to and enriched by another.
- Implementing Retrieval-Augmented Generation (RAG) by embedding and retrieving context from vector databases, then injecting it into AI prompts to improve accuracy and relevance.
- Handling long-running AI processes gracefully — including job orchestration, async background workers, and retry mechanisms.
- Designing safeguards for uncontrolled scenarios: managing failure cases from 3rd party APIs and mitigating the randomness/nondeterminism of LLM outputs.
- Leveraging AI tools and workflows to increase team productivity (e.g., AI-assisted code generation, automated QA, internal bots).
- Writing reusable, testable, and efficient code to improve the functionality of our existing systems.
- Strengthening our test coverage with RSpec to build robust and reliable web apps.
- Conducting full product lifecycles, from idea generation to design, implementation, testing, deployment, and maintenance.
- Providing input on technical feasibility, timelines, and potential product trade-offs, working with business divisions.
- Actively engaging with users and stakeholders to understand their needs and translate them into backend and AI-driven improvements.

About You

We're looking for candidates with a strong track record of working on backend technologies of web apps, ideally with exposure to AI/LLM development or a strong desire to learn.

You should have experience with backend languages and frameworks (Node.js, Django, Rails), as well as modern backend tooling and technologies such as:

- Database management (MySQL, PostgreSQL, MongoDB)
- RESTful APIs
- Security compliance
- Cloud technologies (AWS, Google Cloud, Azure)
- Server-side languages (Java, Python, Ruby, or JavaScript)
- Understanding of frontend technologies
- User authentication and authorization between multiple systems, servers, and environments
- Scalable application design principles
- Creating database schemas that represent and support business processes
- Implementing automated testing platforms and unit tests
- Familiarity with LLM APIs, embeddings, vector databases and prompt design best practices

We're not big on credentials, so a Computer Science degree or graduating from a prestigious university isn't something we emphasize. We care about what you can do and how you do it, not how you got here.

While you'll report to a CTO directly, Rakamin is a company where [Managers of One](#) thrive. We're quick to trust that you can do it, and here to support you. You can expect to be counted on and do your best work and build a career here.

This is a remote job. You're free to work where you work best: home office, co-working space, coffee shops. To ensure time zone overlap with our current team and maintain well communication, we're only looking for people based in Indonesia.

Benefits & Perks

Our benefits support a life well-lived away from work. Ample time off and all the resources you need to support you in doing the best work of your career:

1. **Paid time off:** Rakamin offers **17 days** of vacation and personal days. You can take a day off without any specific reason, whether you want to hang out, take a vacation, or just lay in bed.
2. Learning benefits with total **Rp29 million per year**:
 - Rp6 million per year to subscribe to courses, buy books, or any resources you need to boost your skills and knowledge.
 - O'Reilly subscription worth Rp8 million per year — We subscribe to O'Reilly for you so you can read the newest books, videos, conferences, and courses related to technology, data, design, business, and soft skills.
 - Access to all our learning module Bootcamps and Short Courses worth Rp15 million per year, so you can always upskill and reskill for a better life and career.
3. **Device ownership:** We provide a Rp7 million budget per year so you can purchase any device to support your productivity.